

Gaussian process SDE models

Roland Hostettler¹

Department of Electrical Engineering
Uppsala University
E-Mail: roland.hostettler@angstrom.uu.se

December 7, 2022

¹Based on joint work with Zheng Zhao, Filip Tronarp, and Simo Särkkä

Stochastic differential equations

A typical stochastic differential equation (SDE):

$$d\mathbf{x}_t = \mathbf{f}(\mathbf{x}_t, \mathbf{u}_t, t) dt + \mathbf{L}(\mathbf{x}_t, \mathbf{u}_t, t) d\boldsymbol{\beta}_t$$

where

- ▶ $\mathbf{x}_t \in \mathbb{R}^{d_x}$ is the state,
- ▶ $\mathbf{u}_t \in \mathbb{R}^{d_u}$ is a deterministic input,
- ▶ $\mathbf{f}(\mathbf{x}_t, \mathbf{u}_t, t)$ is the drift function,
- ▶ $\mathbf{L}(\mathbf{x}_t, \mathbf{u}_t, t)$ is the diffusion function, and
- ▶ $\boldsymbol{\beta}_t$ is Brownian motion with diffusion matrix $\boldsymbol{\Sigma}$.

Modeling and identification using SDEs

Typical identification problems

- ▶ The SDE is given (e.g., from first principles);
- ▶ Data is given;
- ▶ **Objective:** Estimate the parameters in the SDE model.

What if...

...we do not know the SDE itself?

- ▶ Can we still use the SDE approach?
- ▶ How should we model the drift (and diffusion) function(s)?

Modeling and identification using SDEs (2/2)

Requirements

The model should...

- ▶ ...be flexible;
- ▶ ...account for different nonlinearities;
- ▶ ...quantify uncertainty;
- ▶ ...scale w.r.t. data and state dimensions.

A few approaches

- ▶ Parametric models (e.g., polynomials);
- ▶ neural SDEs and deep neural SDEs;
- ▶ Gaussian processes (GPs).

This talk: Gaussian processes!

Basic idea

GP SDE model (drift-only):

$$d\mathbf{x}_t = \mathbf{f}(\mathbf{x}_t, \mathbf{u}_t, t) dt + d\boldsymbol{\beta}_t,$$

where

- Drift function

$$\mathbf{f}(\mathbf{x}_t, \mathbf{u}_t, t) = [f_1(\mathbf{x}_t, \mathbf{u}_t, t) \quad \dots \quad f_L(\mathbf{x}_t, \mathbf{u}_t, t)]^\top,$$

- GP prior:

$$f_l(\mathbf{x}_t, \mathbf{u}_t, t) \sim \mathcal{GP}(m(\mathbf{x}_t, \mathbf{u}_t, t), k(\mathbf{x}_t, \mathbf{u}_t, t, \mathbf{x}'_t, \mathbf{u}'_t, t')).$$


Gaussian processes

From random variables to random processes

- ▶ Assume: $f : \mathbb{R}^{d_x} \mapsto \mathbb{R}$ is a function of $\mathbf{x} \in \mathbb{R}^{d_x}$
- ▶ Let $f_i = f(\mathbf{x}_i)$
- ▶ Then, f is a **Gaussian process** if

$$p(f_i, f_j) = \mathcal{N} \left(\begin{bmatrix} f_i \\ f_j \end{bmatrix}; \begin{bmatrix} m_i \\ m_j \end{bmatrix}, \begin{bmatrix} \sigma_{ii}^2 & \sigma_{ij}^2 \\ \sigma_{ji}^2 & \sigma_{jj}^2 \end{bmatrix} \right)$$

for all \mathbf{x}_i and \mathbf{x}_j



Definition

Definition 1: Gaussian process

A Gaussian process is a random function $f : \mathbb{R}^{d_x} \mapsto \mathbb{R}$ with **mean function**

$$\mathbb{E}\{f(\mathbf{x})\} = m(\mathbf{x})$$

and **covariance function**

$$\text{Cov}\{f(\mathbf{x}), f(\mathbf{x}')\} = k(\mathbf{x}, \mathbf{x}')$$

for which it holds that

$$p(f(\mathbf{x}), f(\mathbf{x}')) = \mathcal{N} \left(\begin{bmatrix} f(\mathbf{x}) \\ f(\mathbf{x}') \end{bmatrix} ; \begin{bmatrix} m(\mathbf{x}) \\ m(\mathbf{x}') \end{bmatrix}, \begin{bmatrix} k(\mathbf{x}, \mathbf{x}) & k(\mathbf{x}, \mathbf{x}') \\ k(\mathbf{x}', \mathbf{x}) & k(\mathbf{x}', \mathbf{x}') \end{bmatrix} \right)$$

for all $\mathbf{x}, \mathbf{x}' \in \mathbb{R}^{d_x}$. We denote f as

$$f(\mathbf{x}) \sim \mathcal{GP}(m(\mathbf{x}), k(\mathbf{x}, \mathbf{x}')) .$$

GP regression (1/2)

- Assume that

$$f(\mathbf{x}) \sim \mathcal{GP}(0, k(\mathbf{x}, \mathbf{x}'))$$

$$y = f(\mathbf{x}) + r$$

with $r \sim N(0, \sigma_r^2)$


- Given $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ and

$$\mathbf{y} = [y_1 \quad y_2 \quad \dots \quad y_N]^\top,$$

we have

$$p(\mathbf{f}) = \mathcal{N}(\mathbf{f}; \mathbf{0}, \mathbf{K}_{1:N, 1:N})$$

$$p(\mathbf{y} \mid \mathbf{f}) = \mathcal{N}(\mathbf{y}; \mathbf{f}, \sigma_r^2 \mathbf{I}_N)$$



GP regression (2/2)

- ▶ Given $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N$ and \mathbf{y} , we have

$$p(\mathbf{f}) = \mathcal{N}(\mathbf{f}; \mathbf{0}, \mathbf{K}_{1:N, 1:N})$$
$$p(\mathbf{y} | \mathbf{f}) = \mathcal{N}(\mathbf{y}; \mathbf{f}, \sigma_r^2 \mathbf{I}_N)$$

- ▶ Joint density:

$$p(\mathbf{f}, \mathbf{y}) = \mathcal{N}\left(\begin{bmatrix} \mathbf{f} \\ \mathbf{y} \end{bmatrix}; \begin{bmatrix} \mathbf{0} \\ \mathbf{0} \end{bmatrix}, \begin{bmatrix} \mathbf{K}_{1:N, 1:N} & \mathbf{K}_{1:N, 1:N} \\ \mathbf{K}_{1:N, 1:N} & \mathbf{K}_{1:N, 1:N} + \sigma_r^2 \mathbf{I}_N \end{bmatrix}\right)$$

- ▶ Conditioning on the **measurements** \mathbf{y} yields the **posterior**

$$p(\mathbf{f} | \mathbf{y}) = \mathcal{N}(\mathbf{f}; \mathbf{m}_{\mathbf{f}|\mathbf{y}}, \mathbf{P}_{\mathbf{f}|\mathbf{y}})$$
$$\mathbf{m}_{\mathbf{f}|\mathbf{y}} = \mathbf{K}_{1:N, 1:N} (\mathbf{K}_{1:N, 1:N} + \sigma_r^2 \mathbf{I}_N)^{-1} \mathbf{y}$$
$$\mathbf{P}_{\mathbf{f}|\mathbf{y}} = \mathbf{K}_{1:N, 1:N} - \mathbf{K}_{1:N, 1:N} (\mathbf{K}_{1:N, 1:N} + \sigma_r^2 \mathbf{I}_N)^{-1} \mathbf{K}_{1:N, 1:N}$$

State-space Gaussian process drift model

GP drift model (1)

Assume the scalar, autonomous, time-invariant model

$$\begin{aligned} dx_t &= f(x_t) dt + d\beta_t, \\ f(x_t) &\sim \mathcal{GP}(m(x_t), k(x_t, x'_t)). \end{aligned}$$

Assumptions

- ▶ We perfectly observe a full length N trajectory $x_{0:N} = \{x_0, x_1, \dots, x_N\}$;
- ▶ Mean function is zero, $m(x_t) = 0$;
- ▶ $k(x_t, x'_t)$ is stationary, $k(x_t, x'_t) = k(\Delta x)$ ($\Delta x = x'_t - x_t$)

Sub-optimal solution (1/2)

Euler–Maruyama discretization of the model:

$$x_n \approx x_{n-1} + f(x_{n-1})\Delta t + v_n$$

$$v_n \sim \mathcal{N}(0, \Delta t \sigma^2)$$

► Define $y_n = x_n - x_{n-1}$, then:

$$y_n = f(x_{n-1})\Delta t + v_n$$

$$v_n \sim \mathcal{N}(0, \Delta t \sigma^2)$$

► GP regression:

$$p(\mathbf{f}_{1:N} \mid \mathbf{y}_{1:N}) = \mathcal{N}(\mathbf{f}_{1:N}; \mathbf{m}_{\mathbf{f}|\mathbf{y}}, \mathbf{P}_{\mathbf{f}|\mathbf{y}})$$

$$\mathbf{m}_{\mathbf{f}|\mathbf{y}} = \frac{1}{\Delta t} \mathbf{K}_{1:N,1:N} \left(\mathbf{K}_{1:N,1:N} + \frac{1}{\Delta t} \sigma^2 \mathbf{I}_N \right)^{-1} \mathbf{y}_{1:N}$$

$$\mathbf{P}_{\mathbf{f}|\mathbf{y}} = \mathbf{K}_{1:N,1:N} - \mathbf{K}_{1:N,1:N} \left(\mathbf{K}_{1:N,1:N} + \frac{1}{\Delta t} \sigma^2 \mathbf{I}_N \right)^{-1} \mathbf{K}_{1:N,1:N}^T$$

Sub-optimal solution (2/2)

Solution:

$$y_n = f(x_{n-1})\Delta t + v_n$$

$$v_n \sim \mathcal{N}(0, \Delta t \sigma^2)$$

$$p(\mathbf{f}_{1:n} \mid \mathbf{y}_{1:n}) = \mathcal{N}(\mathbf{f}_{1:n}; \mathbf{m}_{\mathbf{f}|\mathbf{y}}, \mathbf{P}_{\mathbf{f}|\mathbf{y}})$$

with $y_n = x_n - x_{n-1}$.

Drawbacks

- ▶ Euler-Maruyama requires small Δt
 - ▶ ...and higher order discretization schemes require gradient information;
- ▶ Regression scales according to $\mathcal{O}(N^3)$
 - ▶ ...problematic for large N .

Spectral decomposition

Recall:

$$f(x_t) \sim \mathcal{GP}(0, k(x_t, x'_t)).$$

- ▶ The covariance function is assumed **stationary**, $k(x_t, x'_t) = k(\Delta x)$
- ▶ Power spectral density:

$$S_f(\omega) = \mathcal{F}_{\Delta x}\{k(\Delta x)\};$$

- ▶ Decomposition (exact or approximate):

$$S_f(\omega) = q_f H(i\omega) H^*(i\omega);$$

- ▶ $H(i\omega)$ is a **linear system** \Rightarrow state-space representation:

$$dz_x = \mathbf{A}z_x dx + \mathbf{B} d\varepsilon_x,$$

$$f(x) = \mathbf{C}z_x.$$

Resulting strategy

Thus, we can:

1. Sort the state sequence $x_{0:N}$ in ascending order (pseudo-time) $\Rightarrow \{\tilde{x}_{0:N}, \tilde{y}_{1:N}\}$;
2. Consider the model

$$\begin{aligned}\mathbf{z}_n &= \mathbf{F}\mathbf{z}_{n-1} + \mathbf{w}_n, \\ \tilde{y}_n &= \Delta t \mathbf{C} \mathbf{z}_{n-1} + v_n, \\ v_n &\sim \mathcal{N}(0, \Delta t \sigma^2)\end{aligned}$$

3. Run a Kalman filter & Rauch–Tung–Striebel smoother to estimate $p(\mathbf{z}_n \mid \tilde{y}_{1:N})$ and thus $p(f_n \mid \tilde{y}_{1:N})$

Pros and cons

- ▶ Inference scales according to $\mathcal{O}(N)$ (can be implemented as $\mathcal{O}(\log(N))$);
- ▶ Other discretization schemes: Gradients available, but needs xKF/PF/...;
- ▶ Requires sorting the data (incurs a cost).

Example: Double well

- SDE:


$$dx_t = 3(x_t - x_t^3) dt + d\beta_t$$

with $x_0 = 1$;

- Kernel: Matérn,

$$k(\Delta x) = \sigma_M^2 \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\frac{\sqrt{2\nu}\Delta x}{\ell} \right) K_\nu \left(\frac{\sqrt{2\nu}\Delta x}{\ell} \right)$$

with $\nu = 5/2$, $\ell = 1.2$, and $\sigma_M = 3$.



Example: Beneš SDEs

- SDE:



$$dx_t = \tanh(x_t) dt + 0.01 d\beta_t$$

with $x_0 = 0$.

- Kernel: Matérn,

$$k(\Delta x) = \sigma_M^2 \frac{2^{1-\nu}}{\Gamma(\nu)} \left(\frac{\sqrt{2\nu}\Delta x}{\ell} \right) K_\nu \left(\frac{\sqrt{2\nu}\Delta x}{\ell} \right)$$

with $\nu = 5/2$, $\ell = 1.5$, and $\sigma_M = 0.3$.



Reduced-rank Gaussian process drift model

GP drift model (2)

How about the more general, d_x -dimensional model

$$d\mathbf{x}_t = \mathbf{f}(\mathbf{x}_t, \mathbf{u}_t, t) dt + d\boldsymbol{\beta}_t,$$

$$\mathbf{y}_n = \mathbf{g}(\mathbf{x}_{t_n}, \mathbf{u}_{t_n}) + \mathbf{r}_n,$$

with noisy observations $\mathbf{y}_n = \mathbf{y}(t_n)$ and

$$\mathbf{f}(\mathbf{x}_t, \mathbf{u}_t, t) = [f_1(\mathbf{x}_t, \mathbf{u}_t, t) \quad \dots \quad f_L(\mathbf{x}_t, \mathbf{u}_t, t)]^\top,$$
$$f_l(\mathbf{x}_t, \mathbf{u}_t, t) \sim \mathcal{GP}(m(\mathbf{x}_t, \mathbf{u}_t, t), k(\mathbf{x}_t, \mathbf{u}_t, t, \mathbf{x}'_t, \mathbf{u}'_t, t'))?$$

Assumptions

- ▶ Mean function is zero, $m(\mathbf{x}_t, \mathbf{u}_t, t) = 0$;
- ▶ Covariance function is separable,

$$k(\mathbf{x}_t, \mathbf{u}_t, t, \mathbf{x}'_t, \mathbf{u}'_t, t') = k_S(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}'_t, \mathbf{u}'_t) k_T(t, t');$$

- ▶ $k_T(t, t')$ is stationary, $k_T(t, t') = k_T(\tau)$.

Trouble on the horizon...

Model:

$$\begin{aligned}f_l(\mathbf{x}_t, \mathbf{u}_t, t) &\sim \mathcal{GP}(0, k_S(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}'_t, \mathbf{u}'_t)k_T(t, t')), \\dx_{l,t} &= f_l(\mathbf{x}_t, \mathbf{u}_t, t) dt + d\beta_{l,t}, \\ \mathbf{y}_n &= \mathbf{g}(\mathbf{x}_{t_n}, \mathbf{u}_{t_n}) + \mathbf{r}_n.\end{aligned}$$

Problems

- ▶ \mathbf{x}_t is a **latent** state, we only have access to \mathbf{y}_n ;
- ▶ f_l depends on whole trajectories \mathbf{x}_τ and \mathbf{u}_τ for all $\tau \in [0, t]$;
- ▶ **The model is non-Markovian.**

Some solutions

- ▶ Sparse approximations (inducing points);
- ▶ **Basis function expansion;**
- ▶ ...

Basis function expansion (1/2)

Decompose

$$f_l(\mathbf{x}_t, \mathbf{u}_t, t) \sim \mathcal{GP}(0, k_S(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}'_t, \mathbf{u}'_t)k_T(t, t')).$$

in $\mathbf{x}_t, \mathbf{u}_t$ such that

$$f_l(\mathbf{x}_t, \mathbf{u}_t, t) = \sum_{j=0}^{\infty} \alpha_{j,t} \psi_j(\mathbf{x}_t, \mathbf{u}_t),$$

where $\psi_j(\mathbf{x}_t, \mathbf{u}_t)$ are the eigenfunctions of $k_S(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}'_t, \mathbf{u}'_t)$:

$$\langle \psi_i(\mathbf{x}_t, \mathbf{u}_t), \psi_j(\mathbf{x}_t, \mathbf{u}_t) \rangle = \begin{cases} 1 & i = j, \\ 0 & i \neq j, \end{cases}$$

$$\langle k_S(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}'_t, \mathbf{u}'_t), \psi_j(\mathbf{x}'_t, \mathbf{u}'_t) \rangle = \lambda_j \psi_j(\mathbf{x}_t, \mathbf{u}_t),$$

with

$$\langle g(\mathbf{x}), h(\mathbf{x}) \rangle = \int g(\mathbf{x}) h^*(\mathbf{x}) d\mathbf{x}.$$

Basis function expansion (2/2)

Basis function expansion of $f_l(\mathbf{x}_t, \mathbf{u}_t, t)$ in $\mathbf{x}_t, \mathbf{u}_t$:

$$f_l(\mathbf{x}_t, \mathbf{u}_t, t) = \sum_{j=0}^{\infty} \alpha_{j,t} \psi_j(\mathbf{x}_t, \mathbf{u}_t).$$

The (time-varying) coefficients are

$$\alpha_{j,t} = \langle f(\mathbf{x}_t, \mathbf{u}_t, t), \psi_j(\mathbf{x}_t, \mathbf{u}_t) \rangle.$$

► Then:

$$\begin{aligned} \text{Cov}\{\alpha_{i,t}, \alpha_{j,t'}\} &= k_T(t, t') \lambda_j \delta_{ij} \\ \alpha_{j,t} &\sim \mathcal{GP}(0, k_T(t, t') \lambda_j) \end{aligned}$$

► $\alpha_{j,t}$ s are still non-Markovian

Spectral decomposition

Coefficients:

$$\alpha_{j,t} \sim \mathcal{GP}(0, k_{\alpha,j}(t, t'))$$

with $k_{\alpha,j}(t, t') = k_T(t, t')\lambda_j = k_T(\tau)\lambda_j$

- Power spectral density:

$$S_{\alpha,j}(\omega) = \mathcal{F}_{\tau}\{k_{\alpha,j}(\tau)\};$$

- Decomposition:

$$S_{\alpha,j}(\omega) = q_{\alpha,j}H(\mathrm{i}\omega)H^*(\mathrm{i}\omega);$$

- $H(\mathrm{i}\omega)$ is a linear system \Rightarrow state-space representation:

$$\mathrm{d}\mathbf{z}_{j,t} = \mathbf{A}_j\mathbf{z}_{j,t} \mathrm{d}t + \mathbf{B}_j \mathrm{d}\varepsilon_{j,t},$$

$$\alpha_{j,t} = \mathbf{C}_j\mathbf{z}_{j,t}.$$

Dynamic model: Summary (1/2)

- Basis function expansion with finite approximation:

$$\begin{aligned}f_l(\mathbf{x}_t, \mathbf{u}_t, t) &= \sum_{j=0}^{\infty} \alpha_{j,t} \psi_j(\mathbf{x}_t, \mathbf{u}_t) \\ &\approx \boldsymbol{\Psi}(\mathbf{x}_t, \mathbf{u}_t) \boldsymbol{\alpha}_{l,t},\end{aligned}$$

with

- $\boldsymbol{\Psi}(\mathbf{x}_t, \mathbf{u}_t) = [\psi_1(\mathbf{x}_t, \mathbf{u}_t) \quad \psi_2(\mathbf{x}_t, \mathbf{u}_t) \quad \dots \quad \psi_J(\mathbf{x}_t, \mathbf{u}_t)],$
- $\boldsymbol{\alpha}_{l,t} = [\alpha_{1,t} \quad \alpha_{2,t} \quad \dots \quad \alpha_{J,t}]^T;$
- Spectral decomposition:

$$\begin{aligned}d\mathbf{z}_{j,t} &= \mathbf{A}_j \mathbf{z}_{j,t} dt + \mathbf{B}_j d\varepsilon_{j,t}, \\ \alpha_{j,t} &= \mathbf{C}_j \mathbf{z}_{j,t}.\end{aligned}$$

Dynamic model: Summary (2/2)

Complete model for $f_l(\mathbf{x}_t, \mathbf{u}_t, t)$:

$$d\mathbf{z}_{j,t} = \mathbf{A}_j \mathbf{z}_{j,t} dt + \mathbf{B}_j d\varepsilon_{j,t},$$

$$\alpha_{j,t} = \mathbf{C}_j \mathbf{z}_{j,t},$$

$$dx_l = \Psi(\mathbf{x}_t, \mathbf{u}_t) \alpha_{l,t} dt + d\beta_{l,t}.$$


Observations

- ▶ (Partially linear) nonlinear state-space model;
- ▶ The full model is obtained by further stacking;

Encoding structure

- ▶ Prior knowledge of (physical) structure can be encoded;
- ▶ Example: Coupled system:

$$\begin{aligned}dx_{1,t} &= f(x_{1,t}, u_t) dt + d\beta_{1,t}, \\dx_{2,t} &= f(x_{1,t}, x_{2,t}) dt + d\beta_{2,t}.\end{aligned}$$



Estimation

Recall:

$$d\mathbf{z}_{j,t} = \mathbf{A}_j \mathbf{z}_{j,t} dt + \mathbf{B}_j d\varepsilon_{j,t},$$

$$\alpha_{j,t} = \mathbf{C}_j \mathbf{z}_{j,t},$$

$$dx_l = \Psi(\mathbf{x}_t, \mathbf{u}_t) \alpha_{l,t} dt + d\beta_{l,t},$$

$$\mathbf{y}_n = \mathbf{g}(\mathbf{x}_{t_n}, \mathbf{u}_{t_n}) + \mathbf{r}_n$$

Objectives

1. Estimate the Gaussian process' hyperparameters (found in $\mathbf{A}_j, \mathbf{B}_j, \mathbf{C}_j \dots$);
2. Estimate the state \mathbf{x}_t itself;
3. Predict the future state.

Standard tools such as Kalman filters/smoothers
or particle methods can be used.

Examples: Setup

- ▶ Two benchmark problems:
 - ▶ Bouc–Wen oscillator;
 - ▶ Cascaded tanks;
- ▶ Assumptions:
 - ▶ Linear observation models;
 - ▶ Euler–Maruyama discretization;
 - ▶ Fourier basis functions:

$$\psi_j(x) = \frac{1}{\sqrt{\gamma}} \exp\left(\frac{\mathrm{i} j 2\pi x}{\gamma}\right)$$

- ▶ Estimation:
 - ▶ Extended Kalman filter;
 - ▶ Hyperparameters: Maximization of the marginal likelihood $p(\boldsymbol{\theta} \mid \mathbf{y}_{1:N})$;
- ▶ Performance criterion: One-step ahead prediction error.

Example: Bouc–Wen oscillator — Model and data

Model

$$\begin{aligned}dx_{1,t} &= x_{2,t} dt, \\dx_{2,t} &= f(x_{1,t}, x_{2,t}, u_t, t) dt + d\beta_t;\end{aligned}$$

- Covariance functions:

$$\begin{aligned}k_S(\mathbf{x}_t, u_t, \mathbf{x}'_t, u'_t) &= k_{SE}(\mathbf{x}_t, \mathbf{x}'_t) + k_{SE}(u_t, u'_t) \\k_T(t, t') &= k_{OU}(t, t');\end{aligned}$$

- Truncate at 25 (\mathbf{x}_t) and 5 (u_t) eigenfunctions.



Data

- Separate training and validation datasets;
- Two excitations: Swept sine and multi-sine.

Example: Bouc–Wen oscillator — Results (1/2)

Prediction RMSE:

- ▶ Training: 2.65×10^{-5}
- ▶ Multisine validation: 0.580×10^{-5}
- ▶ Swept sine validation: 0.096×10^{-5}



Example: Bouc–Wen oscillator — Results (2/2)

Comparison of the RMSE and R^2 values for the Bouc–Wen benchmark.

Model	Multisine		Swept sine	
	RMSE	R^2	RMSE	R^2
AR(1)	21.6×10^{-5}	0.676	20.9×10^{-5}	0.685
BLA ¹	1.13×10^{-5}	0.983	0.698×10^{-5}	0.989
Volterra ¹	0.895×10^{-5}	0.986	0.347×10^{-5}	0.994
Proposed	0.580×10^{-5}	0.991	0.096×10^{-5}	0.998

¹Schoukens and Griesing-Scheiwe (2016)

Example: Cascaded tanks — Model and data

Model

$$\begin{aligned}dx_{1,t} &= f_1(x_{1,t}, u_t, t) dt + d\beta_{1,t} \\dx_{2,t} &= f_2(x_{1,t}, x_{2,t}, t) dt + d\beta_{2,t}\end{aligned}$$

- Covariance function:

$$k(\mathbf{x}, \mathbf{x}', t, t') = k_{\text{SE}}(\mathbf{x}, \mathbf{x}')k_{\text{OU}}(t, t').$$

- Remaining parameters: Same as for Bouc–Wen



Data

- Separate training and validation datasets;
- Arbitrary operational data.

Example: Cascaded tanks — Results (1/2)

Prediction RMSE:

- ▶ Training set RMSE: 51.5×10^{-3}
- ▶ Validation set RMSE: 57.6×10^{-3} .



Example: Cascaded tanks — Results (2/2)

Comparison of the RMSE and R^2 values for the cascaded tanks benchmark.

Model	RMSE	R^2
AR(1)	185.9×10^{-3}	0.911
BLA ¹	55.6×10^{-3}	0.974
Volterra ¹	49.4×10^{-3}	0.991
GP drift	57.6×10^{-3}	0.972

¹Schoukens and Griesing-Scheiwe (2016)

Some pitfalls

- ▶ Truncation order J depends on the hyperparameters \Rightarrow need to be chosen carefully and taken into account when estimating the hyperparameters;
- ▶ Fourier basis functions are **local**:

$$\psi_j(x) = \frac{1}{\sqrt{\gamma}} \exp\left(\frac{i j 2\pi x}{\gamma}\right)$$

- ▶ γ controls the size of the neighborhood;
 - ▶ Large $\gamma \Rightarrow$ large J ;
- ▶ Large $J \Rightarrow$ high-dimensional model (but conditionally linear in most states):

$$d\mathbf{z}_{j,t} = \mathbf{A}_j \mathbf{z}_{j,t} dt + \mathbf{B}_j d\varepsilon_{j,t},$$

$$\alpha_{j,t} = \mathbf{C}_j \mathbf{z}_{j,t},$$

$$dx_l = \Psi(\mathbf{x}_t, \mathbf{u}_t) \alpha_{l,t} dt + d\beta_{l,t},$$

$$\mathbf{y}_n = \mathbf{g}(\mathbf{x}_{t_n}, \mathbf{u}_{t_n}) + \mathbf{r}_n$$

- ▶ Poor scaling of higher-dimensional basis functions that are based on cartesian products of low-dimensional ones.

Summary

State-space GP SDE model

► Model:

$$\begin{aligned}dx_t &= f(x_t) dt + d\beta_t, \\ f(x_t) &\sim \mathcal{GP}(m(x_t), k(x_t, x'_t)).\end{aligned}$$

► Euler–Maruyama discretization:



$$\begin{aligned}y_n &= f(x_{n-1})\Delta t + v_n \\ v_n &\sim \mathcal{N}(0, \Delta t\sigma^2)\end{aligned}$$

with $y_n = x_n - x_{n-1}$.

► Spectral decomposition (in Δx):

$$\begin{aligned}\mathbf{z}_n &= \mathbf{F}\mathbf{z}_{n-1} + \mathbf{w}_n, \\ \tilde{y}_n &= \Delta t\mathbf{C}\mathbf{z}_{n-1} + v_n, \\ v_n &\sim \mathcal{N}(0, \Delta t\sigma^2)\end{aligned}$$

► Estimation using RTS smoother



Reduced-rank GP SDE model

- Proposed model:

$$d\mathbf{x}_t = \mathbf{f}(\mathbf{x}_t, \mathbf{u}_t, t) dt + d\boldsymbol{\beta}_t$$



$$f_l(\mathbf{x}_t, \mathbf{u}_t, t) \sim \mathcal{GP}(0, k_S(\mathbf{x}_t, \mathbf{u}_t, \mathbf{x}'_t, \mathbf{u}'_t)k_T(t, t'))$$

- Basis function expansion and spectral decomposition yield:

$$d\mathbf{z}_{j,t} = \mathbf{A}_j \mathbf{z}_{j,t} dt + \mathbf{B}_j d\varepsilon_{j,t},$$







$$\alpha_{j,t} = \mathbf{C}_j \mathbf{z}_{j,t},$$

$$dx_l = \boldsymbol{\Psi}(\mathbf{x}_t, \mathbf{u}_t) \boldsymbol{\alpha}_{l,t} dt + d\beta_{l,t}$$



References

References

-  C. E. Rasmussen and C. K. I. Williams, *Gaussian Processes for Machine Learning*. The MIT Press, 2006.
-  J. Quiñonero-Candela and C. E. Rasmussen, “A unifying view of sparse approximate Gaussian process regression,” *Journal of Machine Learning Research*, vol. 6, pp. 1939–1959, 2005.
-  S. Särkkä, A. Solin, and J. Hartikainen, “Spatiotemporal learning via infinite-dimensional Bayesian filtering and smoothing: A look at Gaussian process regression through Kalman filtering,” *IEEE Signal Processing Magazine*, vol. 30, no. 4, pp. 51–61, July 2013.
-  A. Solin and S. Särkkä, “Hilbert space methods for reduced-rank Gaussian process regression,” 2014, arXiv:1401.5508.
-  R. Hostettler, F. Tronarp, and S. Särkkä, “Modeling the drift function in stochastic differential equations using reduced rank Gaussian processes,” in *18th IFAC Symposium on System Identification (SYSID)*, Stockholm, Sweden, July 2018.
-  Z. Zhao, F. Tronarp, R. Hostettler, and S. Särkkä, “State-space Gaussian process for drift estimation in stochastic differential equations,” in *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, Barcelona, Spain, May 2020.